

logistic_regression_regularization

January 12, 2017

1 Logistic Regression with Regularization

```
In [1]: import helpers as hlp
import matplotlib.pyplot as plt
import numpy as np
import tensorflow as tf
```

1.1 Configuration

```
In [2]: # Configuration Params
learning_rate = 0.001
epochs = 50000
display_step = 100

degree = 6 # feature mapping degree
lambda = 1000 # regularization parameter
```

1.2 Import Data

```
In [3]: # Import data
train_data = np.loadtxt(
    "../.../data/ex2/ex2data2.txt", dtype="float64", delimiter=",")

train_x_unbiased = np.delete(train_data, -1, axis=1)
train_x = hlp.map_features(
    train_x_unbiased[:, 0], train_x_unbiased[:, 1], degree=degree)

train_y = np.delete(train_data, np.s_[:-1], axis=1)
```

1.3 Visualize Data

```
In [4]: # Visualize Data
labels = {
    "figure_label": "Microchip Test",
    "x_label": "Microchip Test 1",
    "y_label": "Microchip Test 2",
}
```

```
hlp.plot_classified_data(train_x, train_y, **labels)

<generator object macro at 0x105ba8518>
```

1.4 Define Model

```
In [5]: # Parameters
X = tf.placeholder(tf.float64, name="x")
Y = tf.placeholder(tf.float64, name="y")
theta = tf.Variable(
    tf.zeros(
        (train_x.shape[1], 1), dtype=tf.float64), name="theta")
theta_reg = tf.placeholder(tf.float64, name="theta_reg")

m = train_x.shape[0]

# Model
hypothesis = tf.sigmoid(tf.matmul(X, theta))

# Cost function
cost = tf.reduce_sum(1 / m *
    (-Y * tf.log(hypothesis) -
    (1 - Y) * tf.log(1 - hypothesis))) + lambda / (
    2 * m) * tf.reduce_sum(tf.pow(theta_reg, theta_reg))

# Optimizer
optimizer = tf.train.GradientDescentOptimizer(learning_rate).minimize(cost)
```

1.5 Initialization

```
In [6]: # Initialize Session
sess = tf.InteractiveSession()
tf.global_variables_initializer().run()
epoch = 0

file_writer = tf.summary.FileWriter("logs/reg/run1", sess.graph)
```

1.6 Run Model

```
In [7]: # Run Model
for epoch in range(epochs):
    th_reg = theta.eval()[1:]
    sess.run(optimizer, feed_dict={X: train_x, Y: train_y, theta_reg: th_reg})

# Print adaption progress
if epoch % display_step == 0:
    predictions = np.round(hypothesis.eval(feed_dict={X: train_x}))
    accuracy = hlp.get_accuracy(predictions, train_y)
```

1.7 Results

```
In [8]: # Results
        t = theta.eval()
        print("Optimized Weights \n" + str(t) + "\n")
        print("Accuracy: " + str(
            hlp.get_accuracy(
                np.round(hypothesis.eval(feed_dict={X: train_x})), train_y)))
        print("Cost: " + str(
            cost.eval(feed_dict={X: train_x,
                                Y: train_y,
                                theta_reg: th_reg})))
        print("Iterations: " + str(epoch))

        # Decision Boundary
        hlp.plot_decision_boundary(train_x, train_y, t, **labels)
```

Optimized Weights

```
[[ 8.45388093e-01]
 [ 3.03653654e-01]
 [ 8.10667501e-01]
 [-1.37995246e+00]
 [-4.73503164e-01]
 [-7.71688282e-01]
 [-3.93057902e-02]
 [-2.09925289e-01]
 [-2.09569356e-01]
 [-2.04606642e-01]
 [-1.07238150e+00]
 [-4.86305709e-02]
 [-3.91402451e-01]
 [-1.26667411e-01]
 [-8.60193558e-01]
 [-2.66511270e-01]
 [-1.35792011e-01]
 [-5.60339031e-02]
 [-1.71390274e-01]
 [-1.55952692e-01]
 [-5.01793827e-01]
 [-8.02218461e-01]
 [-4.44834226e-03]
 [-1.96750327e-01]
 [-4.62913828e-04]
 [-2.15522237e-01]
 [-5.87175804e-02]
 [-8.13365688e-01]]
```

Accuracy: 82.20338983050848

Cost: nan

Iterations: 49999

<generator object macro at 0x105ba83b8>

1.8 Close Session

```
In [9]: sess.close()
```